



Agentic_AI_Lab1

StudentName- Ravi Deo

System id-2023458249

Roll-2301010710

CSE-F(2023-27)

💡 STEP-BY-STEP WORKING OF THE CODE(lab-1)

◆ 1. Environment Setup

```
!pip install git+https://github.com/huggingface/transformers.git@main
```

```
!pip install -q datasets
```

◆ What this does:

- Installs the **latest Transformers library** from Hugging Face
- Installs the **datasets** library to load ready-made datasets

✗ Required for:

- Loading pre-trained models
 - Loading image-caption datasets
-

◆ 2. Load the Image Captioning Dataset

```
from datasets import load_dataset
```

```
dataset = load_dataset("ybelkada/football-dataset", split="train")
```

◆ What happens here:

- Downloads a **football image dataset**
- Each data item contains:
 - **an image**
 - **a text caption**

Dataset structure:

```
dataset[i]["image"] # image
```

```
dataset[i]["text"] # caption
```

◆ 3. Inspect Dataset Samples

```
dataset[0]["text"]
```

```
dataset[0]["image"]
```

◆ Purpose:

- To verify data
- Ensures images and captions are correctly loaded

☒ This step is mainly for **understanding and debugging**

◆ 4. Create a Custom PyTorch Dataset

```
from torch.utils.data import Dataset, DataLoader
```

A custom class (already defined in notebook):

◆ What this class does:

- Converts **raw images + captions** into model-ready format
- Uses **processor** to:
 - Resize images
 - Tokenize text
 - Convert everything into tensors

☒ Output format:

```
{  
    "pixel_values": image tensor,  
    "input_ids": tokenized caption,  
    "attention_mask": mask  
}
```

This is **required for training with PyTorch**.

◆ 5. Load Processor and Model

```
from transformers import AutoProcessor, BlipForConditionalGeneration
```

```
processor = AutoProcessor.from_pretrained("Salesforce/blip-image-captioning-base")
```

```
model = BlipForConditionalGeneration.from_pretrained("Salesforce/blip-image-captioning-base")
```

◆ Processor:

- Handles **both image + text**

- Prepares input exactly how BLIP expects

◆ Model:

- A pre-trained BLIP image captioning model
 - Already trained on large datasets
 - We fine-tune it on football images
-

◆ 6. Create DataLoader

```
train_dataset = ImageCaptioningDataset(dataset, processor) train_dataloader  
= DataLoader(train_dataset, shuffle=True, batch_size=2)
```

◆ What this does:

- Wraps dataset into batches
- batch_size=2 → trains on 2 images at a time
- shuffle=True → improves learning

☞ Needed for efficient training

◆ 7. Training the Model

```
optimizer = torch.optim.AdamW(model.parameters(), lr=5e-5)
```

◆ Optimizer:

- Updates model weights
 - AdamW is best for transformer models
-

Training Loop

```
for epoch in range(2):
```

```
    for batch in train_dataloader:  
        outputs = model(**batch)  
        loss = outputs.loss  
        loss.backward()  
        optimizer.step()  
        optimizer.zero_grad()
```

◆ What happens internally:

1. Model predicts captions
2. Calculates **loss** (difference from true caption)
3. Backpropagation updates weights
4. Model slowly learns football-specific captions

 This is **fine-tuning**, not training from scratch

◆ 8. Inference (Testing the Model)

```
example = dataset[0]
```

```
image = example["image"]
```

◆ Selects a sample image

◆ Generate Caption

```
inputs = processor(images=image, return_tensors="pt") outputs  
= model.generate(**inputs)  
generated_caption = processor.decode(outputs[0], skip_special_tokens=True) print(generated_caption)
```

◆ What happens:

1. Image is processed
2. Model predicts caption tokens
3. Tokens converted back to human-readable text

 Output:

 Automatically generated image caption

◆ 9. Additional Image Checks

```
dataset[1]["image"]
```

```
dataset[3]["image"]
```

Used to visually verify dataset images.